






next token prediction

prompt = $[x_1 \ x_2 \ x_3]$

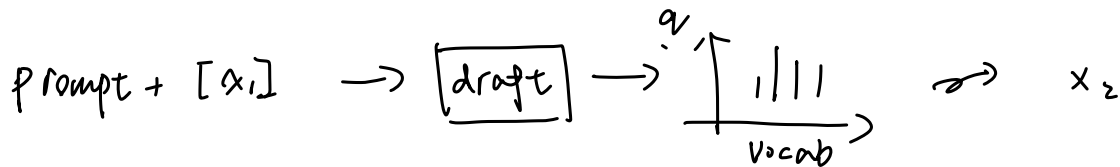
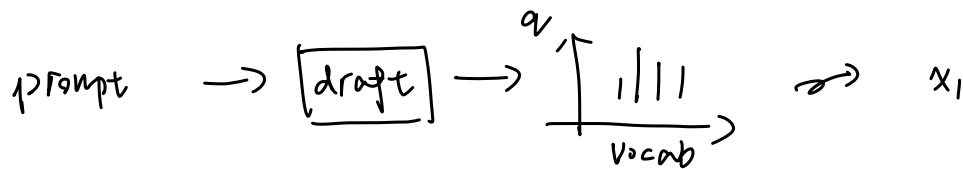


Runtime breakdown: 5% prefix hit, 95% decode

* Running through full decoding is expensive



draft



verify

prompt + $[x_1, x_2, x_3]$



target model

