

The Grader, Stage by Stage

- **1 · TaskTyping [LLM]**. Labels the task as **one of 8 fixed families** (information-retrieval, planning, communication, ...) and picks a scoring regime. A **constrained choice**, not free-form; a deterministic *family* → *regime* rule then overrides the LLM for reliability (and task frontmatter can pre-declare it, skipping the LLM call).
- **2 · TrajectoryNormalizer [det]**. Turns the raw transcript into a clean **Episode**: an ordered step sequence, each `toolResult` attached to its call, per-tool durations recovered from the log.
- **3 · EvaluationPlanner [det]**. **Routing** only, selects the scorer (objective / constraint+utility / judge-based) from the typing.
- **3b · Autograder [det]**. Runs the task's `grade()` checks → `automated_score`, and hands those PASS/FAIL results to the judge as grounding evidence.
- **4 · Regime judge [LLM]**. Scores under the chosen regime with fixed weights; starts at **1.0**, deducts only with a stated reason → judge score.
- **5 · ExecutionTraceAnalyzer [LLM]**. Step-by-step diagnosis: **first-wrong-step · failure-stage · root-cause**, with a per-step score impact. (*Then Latency [det], MDP view [LLM], and the Presenter [det] weighted final score.*)