

Goal + Pipeline Overview

- **Goal. Automatic, reproducible** grading that returns both a **score** and a **diagnosis**.
- **Two evaluators on the same transcript [MINE]**. A fast **deterministic autograder** (runs the task's `grade()` checks → `automated_score`) and a **staged LLM-judge pipeline** (task-typing → normalize → plan → regime judge → trace / MDP → presenter → judge score + diagnosis). Each emits its own score.
- **Task-typed scoring**. A task's family selects a scoring regime with fixed weights (objective = outcome 0.70 / tool-use 0.30; hybrid = outcome 0.30 / constraint 0.30 / preference 0.20 / tool-use 0.20), so scoring is consistent rather than ad hoc.
- **Three runtimes, one grader**. OpenClaw (CLI) · a local mini-agent · KClaw on a real Galaxy S24, all emit the same transcript contract.

The Grader: two evaluators on the same transcript

one input, scored two ways, a fast deterministic autograder and a full LLM-judge pipeline, each emits its own score - [lib_evaluator.py](#)

